

Mandatory Assignment 5

Regular Expressions & Web Scraping (30 + 15 Points/40 + 5 Points)

University of Oslo - IN3110/IN4110

Fall 2020

Your solution to this mandatory assignment needs to be placed in the directory **assignment5** in your Github repository. The repository has to contain a **README.md** file with information on how to run your scripts, required dependencies and packages (versions the code ran with) as well as how to install them. Documentation on how to run your examples¹ is required. Furthermore, your code needs to be well commented and documented. **All** functions need to have docstrings explaining what the function does, how it should be used, an explanation of the parameters and return value(s) (including types). We **highly** recommend you use a well-established docstring style such as the Google style docstrings². However, you are free to choose your own docstring style. We expect your code to be well formatted and readable. Coding style and documentation will be part of the point evaluation for **all tasks** in this assignment.

Files to Deliver for this Task

All **.py**-files will be delivered in the same **assignment5** folder. Each subtask will create an output folder with all outputs as stated in the subtask.

Files required for all tasks

- **README.md** including dependencies and their installation and commands you used for running your examples and creating your output files

Note: Regular expressions can be tricky to write, and they can get hairy very quickly. In addition to being difficult to write, regular expressions are even harder to read. Documentation for your regular expressions should also be given. You are also required to supply files that can be used to demonstrate your solutions.

- You are required to supply everything necessary to demonstrate your code, i.e. you should include the code examples you use when testing your work, or something equivalent.

¹Not only important to get all of the credits, but it is a really important and not understated habit to be established :)

²https://sphinxcontrib-napoleon.readthedocs.io/en/latest/example_google.html

5.1 Sending url requests (2 Points)

In order to parse html and become the regex master of the year, you need to access the html body first. This can be done using the python request module ³.

The module can be installed as follows:

```
pip install requests
```

To retrieve data from any destination, the most common approach is the **get** Request. This is the most used HTTP method. In our example, we will use the get request to fetch data from the website of choice. The result will be printed out as html data. The request module can be imported and used as shown below.

```
#import the module
import requests as req

#grabbing the content of https://en.wikipedia.org/wiki/URL
resp = req.get("https://en.wikipedia.org/wiki/URL")

#get() method returns a response object
print(resp.text)
<!DOCTYPE html>
<html class="client-nojs" lang="en" dir="ltr">
<head>
<meta charset="UTF-8"/>
<title>URL - Wikipedia</title>

#shortened for the purpose of viewing
```

Fantastic, now we can actually display the data in html!

Often you are interested in a specific type of data. With GET requests, you can send parameters in the form of query strings. How does that work? If you were to construct the URL by hand, the data you are interested in would be given as key/value pairs in the URL after a question mark, e.g. `httpbin.org/get?key=val`. Using get requests, you can provide these arguments as a dictionary of strings, using the `params` keyword argument. That means, if you wanted to pass `key1=value1` and `key2=value2` to `httpbin.org/get` the code would in theory look like this:

```
import requests as req

params = {'key1': 'value1', 'key2': 'value2'}
r = req.get('https://httpbin.org/get', params=params)

# the url would be encoded accordingly
print(r.url)
https://httpbin.org/get?key2=value2&key1=value1
```

³<https://requests.readthedocs.io/en/master/user/install/>

For a real example you can look at the following snippet:

```
import requests as req

params = {'user_name': 'admin', 'password': 'password'}
r = req.get('http://httpbin.org/get', params=params)

print(r.url)
http://httpbin.org/get?user_name=admin&password=password

# print the .text content of the request
print(r.text)

{
  "args": {
    "password": "password",
    "user_name": "admin"
  },
  "headers": {
    "Accept": "*/*",
    "Accept-Encoding": "gzip, deflate",
    "Host": "httpbin.org",
    "User-Agent": "python-requests/2.22.0",
    "X-Amzn-Trace-Id": "Root=1-5f841066-0c0ea5084573be8a230feaea"
  },
  "origin": "51.175.234.27",
  "url": "http://httpbin.org/get?user_name=admin&password=password"
}
```

Your task will be to create the function `get.html` which makes a url request of a given website. You should test your function on the following websites:

- https://en.wikipedia.org/wiki/Studio_Ghibli
- https://en.wikipedia.org/wiki/Star_Wars
- https://en.wikipedia.org/wiki/Dungeons_%26_Dragons

Furthermore you should extend the function to optionally take in parameters to be passed to the get function. There should be an optional argument allowing to specify that the response url and text will be saved to a text file with a specified name `optionalargument.txt`. If the optional argument is not set, the response is just returned.

The function should be able to be used like this:

```
get.html(url, params, output)
```

```
# if specified, the response txt and url get printed to a .  
txt — file with the name output.txt
```

The extended function with parameters should be tested for these websites:

- <https://en.wikipedia.org/w/index.php>
 - with parameters title=Main.Page and action=info
- <https://en.wikipedia.org/w/index.php>
 - with parameters title=Hurricane.Gonzalo and oldid=983056166

Files to Deliver in this Subtask

Create a folder `requesting_urls` where you put all output files from our test runs created for solving this subtask.

Files Required in this Subtask

- `requesting_urls.py`
- Five files inside the `requesting_urls` folder containing the function output for the five websites example websites.

5.2 Regex for filtering URLs (5 Points)

In this task, you will be making functions for finding urls in a body of html using regex. Create a script named `filter_urls.py` that includes a function `find_urls` that receives a string of html and returns a list of all urls found in the text. You will be using the `re.findall` method taught in the lecture. This will give us a list of all the strings matching your regular expression. To be strict, we only consider urls that are anchor⁴ hyperlinks with the `<a>` HTML tag. The function should ignore fragment identifiers. That is, links that start with a hash (`#`) symbol. If the url points to a different page, but a specific fragment/section using the hash symbol, the fragment part of the url should be stripped before it is returned by the function.

Example: The string `https://www.example.com/somepage#someidentifier` becomes `https://www.example.com/somepage`

In order to handle relative urls, the function could optionally receive a base url.

Further, make a function `find_articles` that calls `find_urls` and returns only urls to Wikipedia articles. This function should also use regex and **be able to handle any chosen language** (`no.wikipedia.org`, `en.wikipedia.org` etc).

Note: Non wikipedia pages can be ignored. Combine these functions with your functions from previous tasks in order to test them on the following websites:

⁴name of the HTML-element, with the HTML-tag "a"

- https://en.wikipedia.org/wiki/Nobel_Prize
- <https://en.wikipedia.org/wiki/Bundesliga>
- https://en.wikipedia.org/wiki/2019%E2%80%9320_FIS_Alpine_Ski_World_Cup

Note 1: You should only return normal Wikipedia articles, so no special namespace article (or files). While you could exclude just the specific reserved namespaces ⁵, it is sufficient to exclude all articles with a colon. This will exclude some actual articles (for example https://en.wikipedia.org/wiki/Avengers:_Endgame, but we won't deduce points for it).

Note 2: You should only find urls that are in the href attribute of the a tag. It is *not always* the first attribute which means matching just `ja href=` is not enough, but you can assume some things about the use of special html characters since regex cannot parse html alone.

Note 3: You need to include relative urls, best be combined with the base url. Examples of relative urls: `/wiki/Executive_director` and `//en.wikipedia.org/wiki/Wikipedia:Contact_us`

The found urls should be saved to a .txt file, which name can be defined by an optional argument. **Files to Deliver in this Subtask**

Create a folder `filter_urls` where you put all output files from our test runs created for solving this subtask.

Files Required in this Subtask

- `filter_urls.py`
- Three files inside the `filter_urls` folder containing a list of the urls returned for each of the three example websites.

Note: You should not use Beautiful Soup or any HTML parser in this task, only regex.

5.3 Regular Expressions for finding Dates (IN3110 OPTIONAL & IN4110 REQUIRED 10 Points)

In this task, you will be making functions for finding "dates" and "years" in a body of html using regex. Create a script named `collect_dates.py` that includes a function `find_dates` that receives a string of html and returns a list of all dates found in the text in the following format:

```
1) 1998/10/12
2) 1998/11/04
3) 1999/01/13
```

Which date formats do we need to consider? You need to consider the wikipedia standardized formats ⁶. These are the four formats:

⁵i.e. look here <https://en.wikipedia.org/wiki/Wikipedia:Namespace>

⁶<https://en.wikipedia.org/wiki/Template:Date>

DMY: 13 October 2020
MDY: October 13, 2020
YMD: 2020 October 13
ISO: 2020-10-13

Note: You also need to include the case, where the month is abbreviated, e.g. Oct for October.

The dates **returned** should be formatted like this **year/month/day**. The function should be able to find dates even if they are written as strings like **March 12, 2013**. The function should furthermore be able to **recognize, if there is no day**, but only the month and the year. Then it should just return the year and the months. For a list of dates found which have two dates with **year/month/day** and one with **year/month** the returned list could look like this:

```
1) 1998/10/12
2) 1998/11
3) 1999/01/13
```

We also want to mention that there is a regional ambiguity with regards to date formatting. For instance, 9/10 2011, could be either October, 9th or September, 10th. You can assume one of them, but mention which one you chose in your **README.md**. Moreover, there should be an optional argument allowing to specify that the resulting list will be saved to a text file with a specified name **optionalargument.txt**.

You will be using the **re.findall** method taught in the lecture. This will give us a list of all the strings matching your regular expression.

Here is the list of websites, you need to deliver a report in form of a text-file for:

Steps for creating the script

- visit the following websites
- https://en.wikipedia.org/wiki/Linus_Pauling
- https://en.wikipedia.org/wiki/Rafael_Nadal
- https://en.wikipedia.org/wiki/J._K._Rowling
- https://en.wikipedia.org/wiki/Richard_Feynman
- https://en.wikipedia.org/wiki/Hans_Rosling

Save your implementation to **collect_dates.py**.

Files to Deliver in this Subtask

Create a folder **filter_dates_regex** where you put all output files from our test runs created for solving this subtask.

Files Required in this Subtask

- **collect_dates.py**

- Five files inside the `filter_dates_regex` folder containing a list of found dates in sorted order, for each website given.

Note: You should not use BeautifulSoup or any HTML parser in this task, only regex.

5.4 Making your Life easier with Soup for filtering datetime objects (IN3110 & IN4110 8 Points)

Need to plan your watch parties for the upcoming skiing season? Why not send out all the dates and times to your friends that are as well into (watching) skiing? In this task you will use the Python tool BeautifulSoup to parse data. You will extract the datetime objects representing the event "date", a regular expression fetching the "venue" and the discipline "type".

You can install BeautifulSoup via:

```
pip install beautifulsoup4
```

Note: Make sure that you are running python3 and your pip points to the right python version as well. If you are running into issues we highly recommend setting up a `virtualenv` or a `conda env`.

Before starting with the fun-coding part, it is recommended to take some time to get familiar with the html tags of the website we are extracting information from.

Note: It will really help for this task if you just take a few minutes to familiarize yourself with the page setup.

We recommend to navigate to the website using your favorite browser, right click and 'Inspect'. Get familiar with the elements used. Alternatively, right click on the web page and view the page source. Since we are interested in a table, search for the table class :

Okay, so we would like to access the table of soups now. To do so we need to first grab the HTML from the web page, so that we have something to parse through. Then we will extract the soup table.

```
from bs4 import BeautifulSoup
import requests as req

url = "https://en.wikipedia.org/wiki/List_of_soups"

# request url as we have already learned - yeah!
request = req.get(url)

soup = BeautifulSoup(request.text, "html.parser")

#check the title of the wikipedia page
print(soup.title)

# get the soup table
```

Figure 1: Table class in the html. You get to this view using "View Source Page" mode.

Great, now we need to know how to navigate through the table. To figure out which column holds which information, it is recommended to look at the header. Therefore, we will use the html `<th>` tag, which defines the header cell in an html table. On a sidenote, a html table has two kinds of cells: header cells created with `<th>` elements and data cells created with `<td>` elements.

Write a script `time_planner.py` that first requests the url from `https://en.wikipedia.org/wiki/2019%E2%80%932020_FIS_Alpine_Ski_World_Cup`. The request can be done with a function created in earlier tasks.

Then parses through the html using `Beautiful Soup` and finds the main table in the calendar section (see figure below).

Write a function `extract_events` that extracts the data in the date, venue and discipline column. The function will live in the script `time_planner.py`. Since you would like to gamble with your friends beforehand you will create a betting slip. That slip will contain and a table which columns represent the date, venue and discipline as well as a column for the tip of which athlete or country wins that competition. The empty betting slip should be created auto-

← → https://en.wikipedia.org/wiki/2020-21_FIS_Alpine_Ski_World_Cup

Calendar [edit]

Event Key: DH – Downhill, SL – Slalom, GS – Giant Slalom, SG – Super Giant Slalom, AC – Alpine Combined, PS – Parallel Slalom, PG – Parallel Giant Slalom

#	Event	Date	Venue	Type	Winner	Second	Third	Details
1783	1	18 October 2020	 Soldeu	GS 455				
1784	2	14 November 2020	 Lenç de Zúrs	PG 407				
		28 November 2020	 Lake Louise	DH 455				
		29 November 2020		SL 455				
		4 December 2020		SG 455				North American Tour cancelled due to the coronavirus pandemic; most likely to be rescheduled in Val d'Isère
		5 December 2020	 Beaver Creek	DH 455				
		6 December 2020		GS 455				
1785	3	5 December 2020		GS 455				
1786	4	6 December 2020	 Val d'Isère	GS 455				
1787	5	12 December 2020		DH 457				
1788	6	13 December 2020	 Val d'Isère	SG 214				
1789	7	18 December 2020	 Val Gardena/Gröden	SG 214				
1790	8	19 December 2020		DH 458				
1791	9	20 December 2020	 Alta Badia	GS 455				
1792	10	21 December 2020		SL 458				
1793	11	22 December 2020	 Madonna di Campiglio	SL 458				
1794	12	28 December 2020		DH 458				
1795	13	29 December 2020	 Bormio	SG 330				
1796	14	6 January 2021	 Zagreb	SL 330				
1797	15	8 January 2021		GS 454				
1798	16	9 January 2021	 Adelboden	GS 455				
1799	17	10 January 2021		SL 331				
1800	18	15 January 2021		DH 330				
1801	19	16 January 2021	 Wengen	DH 331				
1802	20	17 January 2021		SL 332				
1803	21	22 January 2021		SG 331				
1804	22	23 January 2021	 Titisee	DH 332				
1805	23	24 January 2021		SL 333				
1806	24	28 January 2021	 Schladming	SL 334				
1807	25	30 January 2021		SL 335				
1808	26	31 January 2021	 Chamonix	SL 336				
1809	27	5 February 2021	 Garmisch-Partenkirchen	SG 337				
1810	28	6 February 2021		DH 337				
FIS Alpine World Ski Championships 2021 (8-21 February)								
1811	29	27 February 2021	 Bansko	GS 456				
1812	30	28 February 2021		GS 457				
1813	31	6 March 2021		DH 456				
1814	32	7 March 2021	 Kvitfjell	SG 338				
1815	33	13 March 2021	 Kranjska Gora	GS 458				
1816	34	14 March 2021		SL 337				
1817	35	17 March 2021		DH 338				
1818	36	18 March 2021	 Cortina d'Ampezzo	SG 234				

Figure 2: Main table in calendar section.

matically by your script and saved to a new file named `betting_slip_empty.md` using Markdown ⁷ formatting.

Check that your script also works when accessing https://en.wikipedia.org/wiki/2020%E2%80%9321_FIS_Alpine_Ski_World_Cup.

Files to Deliver in this Subtask

Create a folder `datetime_filter` where you put all output files from our test runs created for solving this subtask. *Files Required in this Subtask*

- `time_planner.py`
- `datetime_filter/betting_slip_empty.md` - a nicely formatted table for your bets, containing date, venue and discipline, but with an empty "who wins" column.

The resulting betting slip would look like this:

5.5 NBA Player Statistics Season 2019/2020 (IN3110 & IN4110 15 Points)

Go Miami Heat - Whooo! Let's imagine you are actually into basketball. After this unpredictable season, you really want dive into the player statistics. Lucky you! In this task you are going to write a script `fetch_player`

⁷<https://github.com/adam-p/markdown-here/wiki/Markdown-Cheatsheet#tables>

BETTING SLIP

Name: _____

DATE	VENUE	DISCIPLINE	Who Wins?
01/01/2021	nowhere	sitting	

Figure 3: Betting Slip.

statistics.py which visits the "2020 NBA playoffs" website on wikipedia (this one: https://en.wikipedia.org/wiki/2020_NBA_playoffs) and creates some player statistics. As in task 5.4 will use BeautifulSoup for parsing.

As always, you first need to make a url request using your function `get_html`.

Then you want to navigate, using BeautifulSoup to the main table in the "Bracket" section shown below.

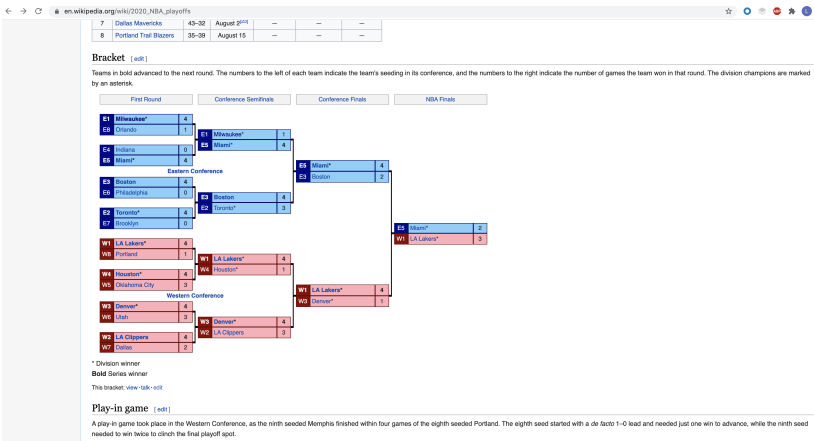


Figure 4: Bracket Section.

There you will extract the names of the teams which made it to the confer-

ence semifinals (This is shown in the section "Bracket").⁸ You will also create a function `extract_url` which will allow you to extract the team urls in this table. You will follow the urls to the wikipedia websites of the corresponding teams in the Bracket section.

⁸Hint: It should be 8.

The next figure shows the website you get to when following the Milwaukee Bucks link in the conference semifinals.

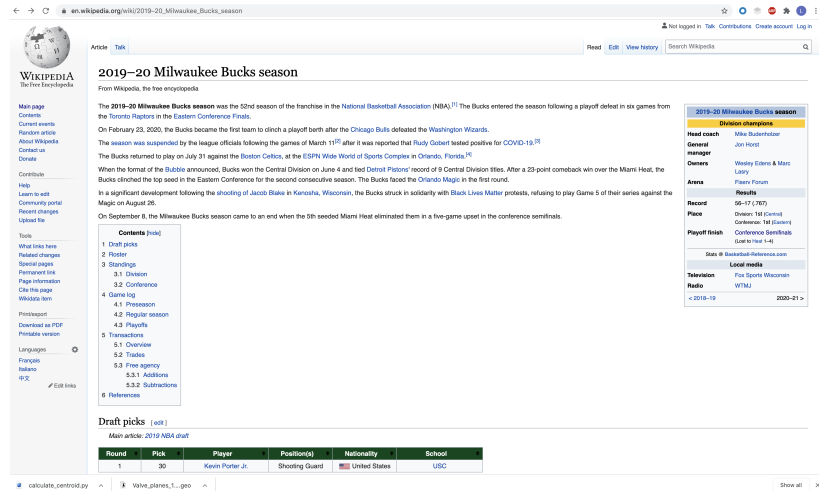


Figure 5: Milwaukee Bucks Page.

In the next step you need to identify all the players that played for the specific team in that season. Therefore you need to navigate to the table holding the team roster. For the Milwaukee Bucks the roster section looks like this:

Roster [edit]

Roster listing

Milwaukee Bucks roster									
Players								Coaches	
Pos.	No.	Name	Height	Weight	DOB (YYYY-MM-DD)	From	Head coach		
F	34	Anelkounmpo, Giannis	6 ft 11 in (2.11 m)	242 lb (110 kg)	1994-12-06	Greece	• Mike Budenholzer		
F	43	Anelkounmpo, Thanasis	6 ft 6 in (1.98 m)	219 lb (99 kg)	1992-07-18	Greece	Assistant coach(es)		
G	6	Belland, Eric	6 ft 9 in (2.06 m)	214 lb (97 kg)	1990-12-09	Kentucky	• Vin Baker		
GMF	23	Brown, Sterling	6 ft 5 in (1.96 m)	219 lb (99 kg)	1995-02-10	SMU	• Chad Foster		
G	24	Connaughton, Pat	6 ft 5 in (1.96 m)	209 lb (95 kg)	1993-01-06	Notre Dame	• Darvin Ham		
G	0	D'Innocenzo, Dante	6 ft 4 in (1.93 m)	203 lb (92 kg)	1997-01-31	Villanova	• Charles Lee		
G	3	Hill, George	6 ft 3 in (1.91 m)	188 lb (85 kg)	1996-05-04	IUPUI	• Josh Longstaff		
F	7	Jayson, Ethan	6 ft 9 in (2.06 m)	235 lb (107 kg)	1997-05-15	Turkey	• Patrick St. Andrews		
GMF	26	Korver, Kyle	6 ft 7 in (2.01 m)	212 lb (96 kg)	1981-03-17	Oregon	• Ben Sullivan		
C	11	Lopez, Brook	7 ft 0 in (2.13 m)	282 lb (128 kg)	1988-04-01	Stanford	Legend		
C	42	Lopez, Robin	7 ft 0 in (2.13 m)	281 lb (127 kg)	1988-04-01	Stanford	• (C) Team captain		
G	15	Mason, Frank (TW)	5 ft 11 in (1.80 m)	190 lb (86 kg)	1994-04-03	Kansas	• (DP) Unsigned draft pick		
G	9	Mathews, Wesley	6 ft 4 in (1.93 m)	220 lb (100 kg)	1986-10-14	Marquette	• (FA) Free agent		
F	23	Maddison, Kris	6 ft 7 in (2.01 m)	222 lb (101 kg)	1991-08-12	Texas A&M	• (S) Suspended		
GMF	13	Reynolds, Cameron (TW)	6 ft 7 in (2.01 m)	225 lb (102 kg)	1995-02-07	Tulane	• (GL) On assignment to G League affiliate		
F	20	Williams, Marvin	6 ft 8 in (2.03 m)	237 lb (108 kg)	1986-06-19	North Carolina	• (TW) Two-way affiliate player		
F	5	Wilson, D. J.	6 ft 10 in (2.08 m)	231 lb (105 kg)	1996-02-19	Michigan	• Injured		
								Roster	
								Last transaction: 2020-02-10	

Standings [edit]


Division [edit]

Central Division	W	L	PCT	GB	Home	Road	Div	GP
x Milwaukee Bucks	56	17	.767	0.0	30-8	26-12	13-1	73
x Indiana Pacers	45	28	.616	11.0	25-11	20-17	8-7	73
Chicago Bulls	22	43	.338	30.0	14-20	8-23	7-9	65
Detroit Pistons	20	46	.303	32.5	11-22	9-24	5-10	66
Cleveland Cavaliers	19	48	.292	33.0	11-25	8-21	4-10	65

Figure 6: Roster Milwaukee Bucks.

Fantastic, let's start on the player statistic section! In order to filter the

statistics for one specific player you need to extract the player name and the url to this players wikipedia page. You will request the `url` to that player which can be found in the roster. If we follow the link to Giannis Antetokounmpo this would be the part of the webpage showing the player statistics.



The screenshot shows the Wikipedia page for Giannis Antetokounmpo. The page is in English (en.wikipedia.org/wiki/Giannis_Antetokounmpo). A legend at the top defines various statistics abbreviations: GP (Games played), GS (Games started), MPG (Minutes per game), FG% (Field goal percentage), 3P% (Three-point field goal percentage), FT% (Free throw percentage), RPG (Rebounds per game), APG (Assists per game), SPG (Steals per game), BPG (Blocks per game), PPG (Points per game), and Bold (Career high). The main table is divided into 'Regular season' and 'Playoffs' sections. The 'Regular season' table has columns for Year, Team, GP, GS, MPG, FG%, 3P%, FT%, RPG, APG, SPG, BPG, and PPG. The 'Playoffs' table has the same columns. Below the tables, there is a section for 'Awards and honours' listing various achievements.

Year	Team	GP	GS	MPG	FG%	3P%	FT%	RPG	APG	SPG	BPG	PPG
2013-14	Milwaukee	77	23	24.6	.414	.347	.683	4.4	1.9	.8	.0	6.8
2014-15	Milwaukee	81	71	31.4	.491	.159	.741	6.7	2.6	.9	1.0	12.7
2015-16	Milwaukee	80	79	35.3	.506	.257	.724	7.7	4.3	1.2	1.4	16.9
2016-17	Milwaukee	80	80	36.6	.525	.272	.770	8.7	5.4	1.6	1.6	22.9
2017-18	Milwaukee	75	75	36.7	.529	.307	.760	10.0	4.6	1.5	1.4	26.9
2018-19	Milwaukee	72	72	32.8	.578	.296	.729	12.5	5.8	1.3	1.5	27.7
2019-20	Milwaukee	65	63	30.4	.553	.304	.653	13.6	5.6	1.0	1.0	26.5
Career		528	483	32.5	.526	.284	.722	8.9	4.3	1.2	1.3	20.1
All-Star		4	4	28.6	.653	.331	.667	8.8	3.0	1.3	1.0	27.3

Year	Team	GP	GS	MPG	FG%	3P%	FT%	RPG	APG	SPG	BPG	PPG
2015	Milwaukee	6	6	33.5	.366	.000	.798	7.0	2.7	.5	1.5	11.5
2017	Milwaukee	6	6	40.6	.336	.400	.545	9.5	4.0	3.3	1.7	24.8
2018	Milwaukee	7	7	40.0	.370	.396	.691	9.6	6.3	1.4	.9	25.7
2019	Milwaukee	15	15	34.3	.492	.327	.637	12.3	4.9	1.1	2.0	25.5
2020	Milwaukee	9	9	30.8	.559	.325	.580	13.8	5.7	.7	.9	26.7
Career		43	43	35.2	.514	.325	.627	11.0	4.8	1.1	1.5	23.7

Awards and honours

- 2x NBA Most Valuable Player: 2019, 2020
- NBA Defensive Player of the Year: 2020
- 4x NBA All-Star: 2017, 2018, 2019, 2020
- NBA Most Improved Player: 2017
- 4x All-NBA Selection
 - All-NBA First Team: 2019, 2020
 - All-NBA Second Team: 2017, 2018

Figure 7: Giannis Antetokounmpo.

We are interested in the NBA Regular season table. Therefore we need to navigate to this one. We want to extract for the row **Year 2019-20** and the columns points per game, blocks per game and rebounds per game.

For your statistics you are going to compare the best players from the teams in the conference semifinals. "Best" is defined by the highest count of points per game in our case.

The three best players of each team make it into our comparison pool⁹

We want to have a look at the best player with respect to points per game, with respect to blocks per game and rebounds per game.

Therefore we will plot the player over the points/blocks/rebounds. You can choose your favorite plotting tool for this. We would like to have players of each team grouped together (this can be color-coded or with a bracket, or another visualization).

Your implementation should be saved to a file called `fetch_player_statistics.py`.

The produced plots should be saved.

Steps for creating the script

- visit https://en.wikipedia.org/wiki/2020_NBA_playoffs
- crawl through the team and player wikipedia websites via the internal urls as described above

⁹Since we take the 3 best players of each team, we will end up comparing the abilities of 24 players in total.

- define the best 3 players of each team in the conference semifinals based on the "Points per game" in 2019/20
- fetch their statistics in the categories points per game, blocks per game and rebounds per game
- create three plots (players of each team grouped together in the plots):
 - Players over points per game
 - Players over blocks per game
 - Players over rebounds per game
- Save the plots!

Files to Deliver in this Subtask Create a folder `NBA_player_statistics` where you put all output files from our test runs created for solving this subtask.

Files Required in this Subtask

- `fetch_player_statistics.py`
- `NBA_player_statistics/players_over_ppg.png`
- `NBA_player_statistics/players_over_bpg.png`
- `NBA_player_statistics/players_over_rpg.png`

5.6 Challenge - Wiki Race with URLs (5 bonus points and a price for the winner!)

Everyone has probably heard of golf. You try to get the ball into the hole with the least amount of hits. Let us play some Wikipedia golf.

Write a script using your previous functions that finds the shortest way (in number of urls to visit) from https://en.wikipedia.org/wiki/Parque_18_de_marzo_de_1938 to https://en.wikipedia.org/wiki/Bill_Mundell using only urls in Wikipedia articles. The script should work with any wikipedia url. The websites given will be your test case. In order to assign a winner, we will evaluate this for 2 undisclosed urls. The most efficient (and correct) script will win a prize :)

You have to use python ;)

Files to Deliver in this Subtask

Create a folder `wiki_race_challenge` where you put all output files from our test runs created for solving this subtask.

Files Required in this Subtask

- `wiki_race_challenge.py`
- `wiki_race_challenge/shortest_way.txt`

Wohoo - You finished another assignment! Congrats!

Update October, 13th 2020

- check Section 5.2 - anchor explanation added, added requirements for urls, identifier example
- check Section 5.3 - added 3 Notes, defined date formats more precisely)

Update October, 14th 2020

- check Section 5.3 - defined that yyyy/mm[/dd] , where only the day optional

Update October, 15th 2020 Q&A

- **For subtask 5.3, are we meant to convert the date with regular expressions, or can we use a standard library?** - Regex should be used everywhere in task 5.3 and 5.3 - so yes, use regex please.
- **Do we need to use a regex to fetch the venue in 5.4?** - No need to fetch it with regex in 5.4